

Jacobian Space from Scratch: Qwen3.6-27B

Anthropic 的 Jacobian lens 主張語言模型的中間層裡有一個 workspace,放著還沒說出口的想法,可以讀出來也可以改。這份 notebook 用 MLX 在 Apple Silicon 上把方法從零實作一次,再把論文的主要實驗搬到 Qwen3.6-27B-4bit 上跑一遍,記錄哪些成立、哪些不成立。

1. Qwen3.6-27B 的架構,走一次 forward
2. logit lens,以及它讀不到的東西
3. Jacobian lens:把 logit lens 假設成 I 的那一項算出來
4. 換個語言問同一件事
5. J-space 本體:稀疏非負分解,以及它佔 residual 多少
6. 換掉 J-space 裡的概念
7. 論文的五個 workspace 面向,逐項驗證

模型放在 `models/Qwen3.6-27B-4bit` ([mlx-community 版](#)), 跑在 `mlx-vm 0.6.4`。


```
In [1]: from mlx_vlm import load
        from rich import print

def patch_qwen3_5():
    """mlx-vm #1548: 0.6.4 misses the +1.0 shift on qwen3_5 RMSNorm
    weights,
    which garbles every output. Must run before load()."""
    import mlx_vlm.models.qwen3_5.qwen3_5 as _q
    keys = (".input_layernorm.weight", ".post_attention_layernorm.weight",
            "model.norm.weight", ".q_norm.weight", ".k_norm.weight")
    if getattr(_q.Model.sanitize, "_patched", False):
        return
    base = _q.sanitize_key
    def sanitize(self, weights):
        shift = any("mtp." in k for k in weights) or any(
            "conv1d.weight" in k and v.shape[-1] != 1 for k, v in
weights.items())
        weights = {k: v for k, v in weights.items() if "mtp." not in k}
        if self.config.text_config.tie_word_embeddings:
            weights.pop("lm_head.weight", None)
        out = {}
        for k, v in weights.items():
            k = base(k)
            if "conv1d.weight" in k and v.shape[-1] != 1:
                v = v.moveaxis(2, 1)
            if shift and any(k.endswith(s) for s in keys) and v.ndim == 1:
                v = v + 1.0
            out[k] = v
        return out
    sanitize._patched = True
    _q.Model.sanitize = sanitize

patch_qwen3_5()
```

```
/Users/happydog/dev/jspace-repro/.venv/lib/python3.13/site-
packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update
jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
```

1. Qwen3.6-27B 的架構



多模態模型,影像和文字各走一條前處理,併成同一串 sequence 之後送進 decoder。

文字這條:tokenizer 把字串切成 token id,再去 embedding 表查出 5120 維向量。

影像這條:切成 16x16 的 patch,每個 patch 過 vision encoder (ViT),再由 PatchMerger 把 2x2 個 patch 併成一個並投影到 5120 維。Qwen 的 vision encoder 接近 SigLIP-2 那一系。

兩邊出來都是 5120 維,併成同一串之後 decoder 不區分哪些向量來自影像、哪些來自文字。最後一個位置的輸出過 `lm_head` 變成 248,320 個候選字的分數,取最高分就是下一個 token。

經典論文: [Attention Is All You Need](#) (Transformer 本體) · [An Image is Worth 16x16 Words](#) (ViT) · [SigLIP 2](#) (vision encoder) · [Qwen2-VL](#) (圖文怎麼併、MRoPE 怎麼標位置)

```
In [2]: model, processor = load("models/Qwen3.6-27B-4bit")
        tokenizer = processor.tokenizer

        llm = model.language_model
        inner = llm.model
        cfg = model.config.text_config
```


64 層由兩種 block 疊成,每 4 層一組。

```
In [3]: for i, layer in enumerate(inner.layers[:4]):
        kind = "GatedDeltaNet (linear attn)" if layer.is_linear else
        "Attention (full attn)"
        print(f"layer {i:2d}    {kind}")
        print(f"    ...      same pattern x 16  ->  {len(inner.layers)} layers
        total")
```

```
layer  0    GatedDeltaNet (linear attn)
```

```
layer 1 GatedDeltaNet (linear attn)
layer 2 GatedDeltaNet (linear attn)
layer 3 Attention (full attn)
...    same pattern x 16 -> 64 layers total
```

兩種 decoder block



兩種 block 的外框一樣,都是在 residual stream 上加東西:

```
h <- h + mixer(RMSNorm(h))
h <- h + FFN(RMSNorm(h))
```

FFN 是 SwiGLU, 5120 升到 17408 再壓回 5120。差別在中間的 mixer 怎麼取用前文。

Full attention block

標準 causal attention。Qwen 的設定是 24 個 Q head、4 個 KV head (GQA, 6 個 Q 共用一組 KV), head_dim 256。另外三個設定: Q 和 K 各自先過 RMSNorm (QK-Norm) 穩住數值; RoPE 只轉前 64 維, 也就是 256 的 25%, 其餘維度不帶位置; 位置用 MRoPE 分三軸編, 影像 patch 才有二維座標可標。

q_proj 的輸出是 $24 \times 256 \times 2 = 12288$ 維, 一半當 Q, 另一半當 output gate, 這是跟標準實作不同的地方:

```
attn = softmax(Q K^T / sqrt(256)) V
out = o_proj(attn * sigmoid(gate))
```

attention 的結果先被 sigmoid gate 逐維縮放,再經 `o_proj` (6144 -> 5120) 寫回 residual stream,所以寫進去多少由這一層自己決定。

GatedDeltaNet block

GatedDeltaNet 不回頭看整串 sequence,而是維護一個固定大小的 state,邊讀邊更新。每個 head 的 state 是一個 `128 x 128` 矩陣 `S` (value 維 x key 維),結構上是一本查找表, `k` 當索引, `v` 當內容。每讀一個 token 做四件事:

```
S <- g * S          # decay: 舊資訊按比例衰減
d <- beta * (v - S k) # delta rule: 新 value 減掉查表結果
S <- S + d k^T       # write: 只把差額寫回去
y <- S q             # read: 用 q 查表,得到這個位置的輸出
```

中間兩步是 delta rule:寫入前先用 `k` 查一次表裡現有的內容 (`S k`),只把差額 `v - S k` 寫回去,所以同一個 key 重複寫入是覆蓋而不是疊加。

力道由兩個 gate 控制,每個 head 各一個純量:

```
g = exp(-exp(A_log) * softplus(a + dt_bias)) # decay
gate, 落在 (0,1)
beta = sigmoid(b)                             # write
gate
```

`a`、`b` 從當前 token 算出來 (`in_proj_a`、`in_proj_b`,5120 -> 48), `A_log`、`dt_bias` 是學來的參數,所以衰減多少、寫入多強都看內容,不是固定值。

進遞迴之前 `q`、`k`、`v` 先過 depthwise causal conv1d (kernel 4) 加 SiLU,讓相鄰幾個 token 混一下,`q` 和 `k` 再各自 normalize 並縮放。出來之後過一個 gated RMSNorm (gate 由 `in_proj_z` 給),最後 `out_proj` (6144 -> 5120) 寫回 residual stream。

Qwen 這邊是 16 個 `q/k` head、48 個 `v` head,head_dim 都是 128,所以一層的 state 是 `48 x 128 x 128`,約 79 萬個數字,與 sequence 長度無關。

經典論文: [RMSNorm](#) · [SwiGLU](#) · [RoPE](#) · [GQA](#) · [Linear Transformers](#) (linear attention 起點) · [Fast Weight Programmers](#) (delta rule 進到這一系) · [Gated Delta Networks](#) (這個 block 本身) · [Mamba-2](#) (同一系的 state space 做法)

兩者的記憶成本差在會不會隨 sequence 長大。

```
In [4]: attn = inner.layers[3].self_attn
        gdn = inner.layers[0].linear_attn

        kv_per_token = 2 * attn.num_key_value_heads * attn.head_dim
        state_size = gdn.num_v_heads * gdn.head_v_dim * gdn.head_k_dim

        print(f"full attention    {attn.num_attention_heads} Q heads /
        {attn.num_key_value_heads} KV heads x {attn.head_dim}")
        print(f"                KV cache grows: {kv_per_token} numbers per token
        per layer")
        print(f"GatedDeltaNet    {gdn.num_v_heads} v heads x {gdn.head_v_dim} x
        {gdn.head_k_dim} state")
        print(f"                fixed: {state_size:,} numbers per layer, any
        sequence length")
        print(f"\nbreak-even at {state_size // kv_per_token} tokens; context limit
        is {cfg.max_position_embeddings:,}")
```

```
full attention    24 Q heads / 4 KV heads x 256
                  KV cache grows: 2048 numbers per token per layer
GatedDeltaNet    48 v heads x 128 x 128 state
                  fixed: 786,432 numbers per layer, any sequence
length
break-even at 384 tokens; context limit is 262,144
```

3:1 的排法讓大部分層用固定大小的壓縮記憶,每 4 層插一次 full attention 精確取用前文。

輸入長什麼樣

文字和影像各跑一次前處理,印出實際的形狀。

```
In [5]: from PIL import Image

        text = "The capital of France is"
        text_ids = tokenizer.encode(text)

        print("TEXT")
        print(f"  input      : {text!r}")
        print(f"  token ids : {text_ids}")
        print(f"  tokens   : {[tokenizer.decode([i]) for i in text_ids]}")

        img = Image.open("assets/qwen3_6_27b_arch_residual.png")
        vision = processor.image_processor(images=img, return_tensors="np")
        pixel_values, grid = vision["pixel_values"], vision["image_grid_thw"]
        t, h, w = grid[0]
        merge = processor.image_processor.merge_size

        print("\nIMAGE")
        print(f"  file      : {img.size[0]}x{img.size[1]} px")
        print(f"  grid t,h,w : {t},{h},{w} -> {t*h*w} patches")
        print(f"  pixel_values: {pixel_values.shape}    1536 = temporal 2 x RGB 3 x
        16 x 16")
        print(f"  after {merge}x{merge} merge: {(h//merge)*(w//merge)} vectors
        into the decoder")
```

TEXT

```
input      : 'The capital of France is'
```

```
token ids : [760, 6511, 314, 9338, 369]
tokens    : ['The', ' capital', ' of', ' France', ' is']
```

IMAGE

```
file      : 1660x2260 px
grid t,h,w : 1,142,104 -> 14768 patches
pixel_values: (14768, 1536) 1536 = temporal 2 x RGB 3 x 16 x 16
after 2x2 merge: 3692 vectors into the decoder
```

英文那句是 5 個 token,一張圖攤開是 14768 個 patch,2x2 併完還有 3692 個向量 進 decoder。影像佔的 context 比文字高三個數量級,merge 那一步就是為了壓這個數字。

接著走一次 forward 確認模型正常。

```
In [6]: import mlx.core as mx

ids = mx.array(tokenizer.encode("The capital of France is"))[None]
logits = llm(ids).logits
next_id = int(mx.argmax(logits[0, -1]))

print(f"logits      : {logits.shape}")
print(f"next token : {tokenizer.decode([next_id])!r}")
```

```
logits      : (1, 5, 248320)
next token : ' Paris'
```

畫圖的設定

後面每一節的結果都畫成圖,圖下面那幾行 print 是圖上那些點的實際數值,方便對照。圖上的文字一律用英文,matplotlib 的預設字型沒有中文字。

顏色固定:灰色是對照或基準線(logit lens、baseline、原本的答案),藍色是 Jacobian lens 這一側,橘色是被介入之後長出來的東西。

In [7]: `import matplotlib.pyplot as plt`

```
INK, LOGIT, JLENS, SWAPPED, MUTED = "#14171a", "#8a9094", "#2f6f8f",
"#c1662f", "#9aa3a8"

plt.rcParams.update({
    "figure.dpi": 130, "savefig.bbox": "tight", "figure.facecolor":
    "white",
    "font.size": 8.5, "axes.titlesize": 9, "axes.labelsize": 8.5,
    "axes.edgecolor": "#c8cccd", "axes.linewidth": 0.8,
    "axes.spines.top": False, "axes.spines.right": False,
    "axes.titlepad": 6, "axes.labelcolor": "#3f464b",
    "xtick.color": "#5a6166", "ytick.color": "#5a6166",
    "xtick.labelsize": 8, "ytick.labelsize": 8,
    "grid.color": "#e8e8ec", "grid.linewidth": 0.7,
    "legend.frameon": False, "legend.fontsize": 8,
    "lines.linewidth": 1.6, "lines.markersize": 4,
})

def panel(ax, title="", xlabel="", ylabel="", ymax=None):
    """One consistent look for every panel: title left, y grid only."""
    if title:
        ax.set_title(title, loc="left", color=INK)
    ax.set_xlabel(xlabel)
    ax.set_ylabel(ylabel)
    ax.grid(axis="y")
    ax.set_axisbelow(True)
    if ymax is not None:
        ax.set_ylim(0, ymax)
    return ax


print("plot style ready")
```

plot style ready

2. Logit lens

可解釋性 (interpretability) 要問的是模型吐出答案的路上,中間發生了什麼。

切入點是 residual stream。每一層不重寫整條向量,而是把自己算的東西加回同一條向量上,所以 64 層的輸出都停在同一個座標系裡,前面寫進去的東西後面每一層都讀得到 ([A Mathematical Framework for Transformer Circuits](#), Elhage et al. 2021)。



既然每層的輸出都在同一個座標系上,就可以拿同一把尺去量每一層。logit lens ([Interpreting GPT: the logit lens](#), nostalgebraist 2020)把模型最後那顆頭借來,本來只讀最終的 h_L ,現在對

每一層的 h_ℓ 都讀一次:

$$\text{LogitLens}(h_\ell) = W_U \cdot \text{norm}(h_\ell)$$

`norm` 是 final RMSNorm, W_U 是 `lm_head`, 讀出來的是「第 ℓ 層時 模型認為下一個字是什麼」。

要抽 residual 得手動一層層跑, 自己準備兩種 mask 和 MRoPE 位置, 才跟官方 forward 一致。

```
In [8]: from mlx_vlm.models.qwen3_5.language import (
        _create_qwen3_5_attention_mask, _create_qwen3_5_ssm_mask)

dtype = inner.norm.weight.dtype
d_model = cfg.hidden_size
n_layers = len(inner.layers)

def encode(text):
    return mx.array(tokenizer.encode(text))[None]

def prep(h):
    fa = _create_qwen3_5_attention_mask(h, None)
    ssm = _create_qwen3_5_ssm_mask(h, None)
    pos = mx.tile(mx.arange(h.shape[1])[None, None, :], (3, 1, 1))
    return fa, ssm, pos

def set_linear_train(mode):
    """eval uses a non-differentiable metal kernel; train uses pure ops
    that mx.vjp
    can backprop through. Needed later for the Jacobian."""
    for layer in inner.layers:
        if layer.is_linear:
            layer.linear_attn.train(mode)

def residuals(ids):
    set_linear_train(False)
    h = inner.embed_tokens(ids)
    fa, ssm, pos = prep(h)
    out = []
    for layer in inner.layers:
        h = layer(h, mask=(ssm if layer.is_linear else fa),
                    cache=None, position_ids=pos, position_embeddings=None)
        out.append(h)
    return out

def readout(h):
    return llm.lm_head(inner.norm(h))

res = residuals(encode("The capital of France is"))
print(f"{len(res)} residuals, each {tuple(res[0].shape)} = (batch, tokens,
d_model)")
```

64 residuals, each (1, 5, 5120) = (batch, tokens, d_model)

```

In [9]: prompt = "The capital of France is"
ids = encode(prompt)
res = residuals(ids)
last = ids.shape[1] - 1
answer = int(mx.argmax(readout(res[-1])[0, last]))
ANS = tokenizer.decode([answer])


logit_p, tops = [], []
for l in range(n_layers):
    sc = readout(res[l])[0, last].astype(mx.float32)
    logit_p.append(float(mx.softmax(sc)[answer]))
    tops.append(int(mx.argmax(sc)))

hit = next(l for l in range(n_layers) if tops[l] == answer)

fig, ax = plt.subplots(figsize=(5.6, 2.5))
ax.plot(range(n_layers), logit_p, color=LOGIT)
ax.axvline(hit, color=MUTED, lw=0.9, ls=":")
ax.text(hit - 1, 0.9, f"top-1 becomes {ANS!r} at L{hit}", color=MUTED,
        ha="right", fontsize=8)
panel(ax, f"logit lens: p({ANS.strip()}) at the last position", "layer",
      "p(answer)", 1.0)
plt.show()

print(f"answer {ANS!r}; p(answer) every 4th layer:")
print(" " + " ".join(f"L{l}:{logit_p[l]:.3f}" for l in range(0, n_layers,
4)))

```



```

answer ' Paris'; p(answer) every 4th layer:
  L0:0.000 L4:0.000 L8:0.000 L12:0.000 L16:0.000 L20:0.000 L24:0.000
L28:0.000 L32:0.000 L36:0.000 L40:0.000
L44:0.000 L48:0.000 L52:0.000 L56:0.225 L60:0.858

```

前 50 幾層讀不到答案。top-1 是 ' '、'...'、'____'、'陷'、'zl' 這類碎片, p(answer) 都是 0.000。top-1 到 L55 才變成 ' Paris', 而機率要更晚才起來, L56 是 0.225, L60 是 0.858。

讀不到不代表模型內部沒有答案,也可能是這把尺不夠準。把 h_ℓ 丟進最後那顆頭, 隱含的假設是從第 ℓ 層到終點的那幾十層什麼都不做:

$$\text{LogitLens}(h_\ell) = W_U \cdot \text{norm}(h_\ell) \quad \Longleftrightarrow \quad \frac{\partial h_L}{\partial h_\ell} \approx I$$

這個假設對最後幾層還堪用,對中層太粗。答案可能早就算出來了,只是還沒轉到最後那顆頭 讀得到的方向上。

3. Jacobian lens

Anthropic 的做法是把 logit lens 假設成 I 的那一項算出來:

$$J_\ell = \frac{\partial h_L}{\partial h_\ell} \quad \text{JLens}(h_\ell) = W_U \cdot \text{norm}(J_\ell h_\ell)$$

J_ℓ 是一個 5120×5120 的矩陣,描述第 ℓ 層的 residual 動一點、最終的 residual 會怎麼跟著動。它是在某個輸入、某個位置附近展開的一階近似,不是全域的座標轉換。乘上它等於把 h_ℓ 搬到最終 residual 附近再用同一顆頭讀,跟 logit lens 的差別 只有這一步。

官方實作 [anthropics/jacobian-lens](#), 社群解說 [J-Space: Yet Another LLM Mind Reader?](#), 獨立重現與批評 [A Review of Anthropic's Global Workspace Paper](#)。

從零算 J_ℓ

用 `mx.vjp` 反傳:在第 `src` 層之後注入一個 delta,看第 `tgt` 層的 residual 怎麼變。兩個實務細節:

1. GatedDeltaNet 在 eval 模式走的是不可微的 metal kernel,算之前要切到 train 模式 (純 ops 的 chunked scan),算完切回來。
2. chunked scan 內部會呼叫 `mx.async_eval`,在 `mx.vjp` 裡不合法,暫時換成 `no-op`。

矩陣有 5120 列,全算太慢,這裡算前 64 列。

```

In [10]: import contextlib
import time

SKIP_FIRST = 16  # early positions are BOS and formatting, readout there
is noise
TARGET = n_layers - 1

@contextlib.contextmanager
def no_async_eval():
    orig = mx.async_eval
    mx.async_eval = lambda *a, **k: None
    try:
        yield
    finally:
        mx.async_eval = orig

def forward_with_injection(ids, delta, src, tgt):
    h = inner.embed_tokens(ids)
    fa, ssm, pos = prep(h)
    for l, layer in enumerate(inner.layers):
        h = layer(h, mask=(ssm if layer.is_linear else fa),
                    cache=None, position_ids=pos, position_embeddings=None)
        if l == src:
            h = h + delta
        if l == tgt:
            return h
    return h

def jacobian_rows(ids, src, tgt, n_rows):
    seq = ids.shape[1]
    valid = mx.array(list(range(SKIP_FIRST, seq - 1)), dtype=mx.int32)
    delta = mx.zeros((n_rows, seq, d_model), dtype=dtype)
    onehot = (mx.arange(d_model)[None, :] == mx.arange(n_rows)[:,
None]).astype(mx.float32)
    posmask = (mx.arange(seq)[None, :] == valid[:, None]).any(axis=0)
    cot = (onehot[:, None, :] * posmask[None, :,
None]).astype(mx.float32).astype(dtype)

    set_linear_train(True)
    with no_async_eval():
        (_,), (g,) = mx.vjp(
            lambda d: forward_with_injection(ids, d, src, tgt), (delta,),
(cot,))
        rows = mx.take(g, valid, axis=1).mean(axis=1).astype(mx.float32)
        mx.eval(rows)
    set_linear_train(False)
    return rows

long_prompt = ("The history of the Roman Empire spans several centuries and
includes "
               "many emperors, wars, and cultural achievements that shaped
the ancient world.")
LAYER = 54

t0 = time.perf_counter()

```

```
rows = jacobian_rows(mx.array(tokenizer.encode(long_prompt)[:48])[None],
                      src=LAYER, tgt=TARGET, n_rows=64)
print(f"J[L{LAYER}] first 64 rows: {tuple(rows.shape)} in
{time.perf_counter()-t0:.0f}s")
print(f"diagonal entry J[0,0] = {float(rows[0, 0]):.3f}    row norm =
{float(mx.linalg.norm(rows[0])):.3f}")
```

J[L54] first 64 rows: (64, 5120) in 17s

diagonal entry J[0,0] = 0.949 row norm = 1.136

J[0,0] = 0.949, 不是 1。單看這個對角元素,把 J_ℓ 當成 I 的誤差不大,誤差會不會改變讀出來的字要另外測。

換成公開的 lens

上面是用一條 prompt 算的。Anthropic 的做法是對很多條 prompt 取平均,得到一個不綁特定輸入的 lens。這裡載 Neuronpedia 放出來的擬合結果 ([neuronpedia/jacobian-lens](https://neuronpedia.com/jacobian-lens)), 跟從零算的那 64 列比一下。

```
In [11]: _lens = mx.load("models/Qwen3.6-27B-4bit/jlens.npz")
src_layers = [int(x) for x in _lens["__source_layers__"].tolist()]
n_prompts = int(_lens["__n_prompts__"].tolist()[0])
Jlens = {l: _lens[f"J_{l}"] for l in src_layers}

print(f"public lens: layers {src_layers[0]}..{src_layers[-1]}
({len(src_layers)}), "
      f"fitted on {n_prompts} prompts")

a = rows.reshape(-1)
b = Jlens[LAYER].astype(mx.float32)[:64].reshape(-1)
cos = float((a @ b) / (mx.linalg.norm(a) * mx.linalg.norm(b)))
print(f"cos(our 1-prompt rows, public {n_prompts}-prompt lens) =
{cos:.3f}")

def transport(h, l):
    """Move a layer-l residual into the final-residual basis: J_l @ h."""
    J = Jlens[l]
    return (h.astype(J.dtype) @ J.T).astype(dtype)
```

public lens: layers 0..62 (63), fitted on 1000 prompts

cos(our 1-prompt rows, public 1000-prompt lens) = 0.894

cosine 0.894。比的只有 L54 的前 64 列,單一個數字說不上 J_ℓ 與輸入無關,但一條 prompt 算出來的方向跟一千條擬合的結果大致同向。

兩把尺並排

同一次 forward、同一顆頭,唯一的差別是有沒有乘 J_ℓ 。

```


In [12]: prompt = "The capital of France is"
ids = encode(prompt)
res = residuals(ids)
last = ids.shape[1] - 1
answer = int(mx.argmax(readout(res[-1])[0, last]))
ANS = tokenizer.decode([answer])

LS = [l for l in range(44, n_layers) if l in Jlens]
lg, jc, jtop = [], [], []
for l in LS:
    h = res[l][0, last]
    lg.append(float(mx.softmax(readout(h).astype(mx.float32))[answer]))
    jl = readout(transport(h, l)).astype(mx.float32)
    jc.append(float(mx.softmax(jl)[answer]))
    jtop.append(tokenizer.decode([int(mx.argmax(jl))]))

fig, ax = plt.subplots(figsize=(5.8, 3.0))
ax.plot(LS, lg, color=LOGIT, marker="o")
ax.plot(LS, jc, color=JLENS, marker="o")
ax.text(LS[-1] + 0.4, lg[-1], "logit lens", color=LOGIT, va="center",
        fontsize=8)
ax.text(LS[-1] + 0.4, jc[-1], "Jacobian lens", color=JLENS, va="center",
        fontsize=8)
i55 = LS.index(55)
ax.annotate(f"L55: {lg[i55]:.2f} vs {jc[i55]:.2f}", xy=(55, jc[i55]),
           xytext=(50.2, 0.62), color=INK, fontsize=8,
           arrowprops=dict(arrowstyle="→", color="#c8cccd", lw=0.8))
ax.set_xlim(43.5, 69)
panel(ax, f"same residual, same readout head: p({ANS.strip()})", "layer",
      "p(answer)", 1.0)
plt.show()

print(f"{prompt!r} -> {ANS!r}")
print("  logit      " + " ".join(f"{p:.3f}" for p in lg))
print("  jacobian  " + " ".join(f"{p:.3f}" for p in jc))
print("  jacobian top-1 first becomes the answer at L"
      f"{LS[jtop.index(ANS)]}")

```



```
'The capital of France is' -> ' Paris'
```

```
logit      0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000 0.000  
0.000 0.001 0.088 0.237 0.520 0.922 0.856 0.858  
0.940 0.247
```

```
jacobian 0.000 0.000 0.000 0.000 0.000 0.004 0.001 0.001 0.001  
0.012 0.086 0.895 0.961 0.972 0.990 0.987 0.989  
0.980 0.848
```

jacobian top-1 first becomes the answer at L55

L55 那一行:同一條 residual、同一顆讀出頭,logit lens 給 0.09,Jacobian lens 給 0.90。top-1 也差一層,Jacobian lens 在 L55 就讀到 ' Paris'。

一個必要的對照

公開 lens 在擬合時會跳過每條 prompt 的前 16 個 position, [issue #5](#) 就在講這件事。上面那條 prompt 只有 5 個 token,讀的是 position 4,落在 lens 沒擬合過的區間。所以要在前面墊一段無關的話,把要讀的位置推到 16 以後,看結果還在不在。



```
In [13]: FILLER = ("Here are some notes about geography that provide background context "
                  "for the question that follows this introductory sentence. ")
```

```
def lens_curves(prompt, layers):
    ids = encode(prompt)
    res = residuals(ids)
    last = ids.shape[1] - 1
    ans = int(mx.argmax(readout(res[-1])[0, last].astype(mx.float32)))
    lg, jc = [], []
    for l in layers:
        h = res[l][0, last]
        lg.append(float(mx.softmax(readout(h).astype(mx.float32))[ans]))
        jc.append(float(mx.softmax(readout(transport(h,
l)).astype(mx.float32))[ans]))
    return last, lg, jc
```

```
LS = [l for l in range(44, n_layers) if l in Jlens]
CASES = [("short prompt, outside the fitted range",
          "The capital of France is"),
         ("padded prompt, inside the fitted range",
          FILLER + "The capital of France is")]
```

```
fig, axes = plt.subplots(1, 2, figsize=(7.6, 2.8), sharey=True)
out = []
for ax, (title, prompt) in zip(axes, CASES):
    pos, lg, jc = lens_curves(prompt, LS)
    out.append((title, pos, lg, jc))
    ax.plot(LS, lg, color=LOGIT, marker="o", ms=3)
    ax.plot(LS, jc, color=JLENS, marker="o", ms=3)
    panel(ax, f"{title} (position {pos})", "layer", "", 1.0)
axes[0].set_ylabel("p(answer)")
axes[1].legend(["logit lens", "Jacobian lens"], loc="lower right")
plt.show()
```

```
for title, pos, lg, jc in out:
    print(f"{title} (position {pos})")
    for l, a, b in zip(LS, lg, jc):
        if l in (50, 53, 54, 55, 56, 58):
            print(f"    L{l}    logit {a:.3f}    jacobian {b:.3f}")
```



short prompt, outside the fitted range (position 4)

L50	logit	0.000	jacobian	0.001
L53	logit	0.000	jacobian	0.012
L54	logit	0.001	jacobian	0.086
L55	logit	0.088	jacobian	0.895
L56	logit	0.237	jacobian	0.961
L58	logit	0.922	jacobian	0.990

padded prompt, inside the fitted range (position 23)

L50	logit	0.000	jacobian	0.217
L53	logit	0.001	jacobian	0.765
L54	logit	0.002	jacobian	0.796
L55	logit	0.117	jacobian	0.982
L56	logit	0.321	jacobian	0.979
L58	logit	0.970	jacobian	0.998

落差沒有消失,反而更大。擬合過的位置上,L53 是 logit 0.001 對 Jacobian 0.765, L50 是 0.000 對 0.217。

兩張表合起來看,這條 prompt 的答案 L50 就已經在 residual 裡,logit lens 要到 L58 才讀到 0.97。短 prompt 那版低估了落差,沒有製造落差。

4. 換個語言問同一件事

同一個問題用中文問。

```

In [14]: PARIS = {11751, 57590, 109705}    # ' Paris', 'Paris', and the Chinese
        token
        ZH_PARIS = 109705

def lens_trace(prompt, layers, k=3):
    ids = encode(prompt)
    res = residuals(ids)
    last = ids.shape[1] - 1
    print(f"{prompt!r}    output top-1 = "
          f"{tokenizer.decode([int(mx.argmax(readout(res[-1]))[0,
last]]))!r}")
    print("                jacobian top-3                logit top-3")
    for l in layers:
        if l not in Jlens:
            continue
        h = res[l][0, last]
        j = mx.argsort(-readout(transport(h, l)).astype(mx.float32))[:k]
        g = mx.argsort(-readout(h).astype(mx.float32))[:k]
        js = " ".join(f"{tokenizer.decode([int(t)])!r}" for t in j)
        gs = " ".join(f"{tokenizer.decode([int(t)])!r}" for t in g)
        print(f"  L{l:2d}  {js:40s}  {gs}")


ZH = "法国的首都是"
ids = encode(ZH)
res = residuals(ids)
last = ids.shape[1] - 1
LS = [l for l in range(30, n_layers) if l in Jlens]

curves = {}
for name, tok in (("en", 11751), ("zh", ZH_PARIS)):
    for lens in ("logit", "jacobian"):
        curves[(name, lens)] = []
        for l in LS:
            h = res[l][0, last]
            curves[(name, "logit")].append(
                float(mx.softmax(readout(h).astype(mx.float32))[tok]))
            curves[(name, "jacobian")].append(
                float(mx.softmax(readout(transport(h, l)).astype(mx.float32))
[tok]))

fig, axes = plt.subplots(1, 2, figsize=(7.2, 2.8), sharey=True)
for ax, (name, title) in zip(axes, (("en", "readout of ' Paris' (English
token)"),
                                   ("zh", "readout of the Chinese
token"))):
    ax.plot(LS, curves[(name, "logit")], color=LOGIT, marker="o", ms=3)
    ax.plot(LS, curves[(name, "jacobian")], color=JLENS, marker="o", ms=3)
    panel(ax, title, "layer", "", None)
axes[0].set_ylabel("probability")
axes[0].text(LS[-1], max(curves[("en", "jacobian")]) * 0.9, "Jacobian
lens",
              color=JLENS, ha="right", fontsize=8)
axes[0].text(LS[-1], max(curves[("en", "logit")]) * 1.2 + 0.02, "logit
lens",
              color=LOGIT, ha="right", fontsize=8)
plt.show()

```

```
lens_trace(ZH, range(40, 62, 4))
```



'法国的首都是' output top-1 = '哪个'

jacobian top-3

L40 'Paris' 'city' 'cities'
 L44 '____' 'city' 'cities'
 L48 '____' '____' '____'
 L52 '? ' '____' '____'
 L56 '哪里' '哪个' 'Paris'
 L60 'Paris' '巴黎' 'Paris'

logit top-3

'...' '____' '____'
 '...' '____' '<u'
 '<u' '____' '____'
 'gì' '什么' '哪个'
 '哪里' '哪个' '否'
 '哪个' '哪里' '巴黎'

中文問句,Jacobian lens 在 L40 就把英文的 'Paris' 排第一,同一層的 logit lens 還是 '...' 這類碎片。中文的 '巴黎' 要到 L60 才進 top-3,而模型表面上要說的下一個字是 '哪个',在把句子接完(「法国的首都是哪个城市」)。

這條 prompt 的順序是答案先以英文浮出來,中文晚得多,表面還在造句。方向跟 Anthropic 在 [On the Biology of a Large Language Model](#) 講的多語言迴路一致,不過一條 prompt 證不了「不綁語言的概念空間」,只能說看到的順序 跟那個說法不衝突。

換幾個語言,看 Paris 最早在哪一層進 top-5。

```

In [15]: def first_hit(prompt, targets, topk=5):
    ids = encode(prompt)
    res = residuals(ids)
    last = ids.shape[1] - 1
    fj = fl = None
    for l in range(n_layers):
        if l not in Jlens:
            continue
        h = res[l][0, last]
        j = set(int(t) for t in mx.argsort(-readout(transport(h,
l)).astype(mx.float32))[:topk].tolist()))
        g = set(int(t) for t in mx.argsort(-readout(h).astype(mx.float32))
[:topk].tolist()))
        if fj is None and j & targets:
            fj = l
        if fl is None and g & targets:
            fl = l
    return fj, fl


PROMPTS = [("English", "The capital of France is"),
            ("Chinese", "法国的首都"),
            ("French", "La capitale de la France est"),
            ("Japanese", "フランスの首都は")]

rows = [(lang, p) + first_hit(p, PARIS) for lang, p in PROMPTS]

fig, ax = plt.subplots(figsize=(5.6, 2.4))
for i, (lang, _, fj, fl) in enumerate(rows):
    y = len(rows) - 1 - i
    ax.plot([fj, fl], [y, y], color="#c8cccd", lw=1.4, zorder=1)
    ax.scatter([fj], [y], color=JLENS, zorder=2, s=26)
    ax.scatter([fl], [y], color=LOGIT, zorder=2, s=26)
    ax.text(min(fj, fl) - 1.5, y, f"{abs(fl - fj)}", color=INK,
fontsize=8,
            ha="right", va="center")
ax.set_yticks(range(len(rows)))
ax.set_yticklabels([lang for lang, _, _, _ in reversed(rows)])
ax.set_xlim(14, 66)
ax.set_ylim(-0.6, len(rows) - 0.4)
ax.grid(axis="x")
ax.set_axisbelow(True)
ax.set_title("first layer where Paris enters the top-5", loc="left",
color=INK)
ax.set_xlabel("layer")
ax.scatter([], [], color=JLENS, label="Jacobian lens")
ax.scatter([], [], color=LOGIT, label="logit lens")
ax.legend(loc="lower right")
plt.show()

for lang, p, fj, fl in rows:
    print(f"{lang:9s} {p!r:34s} jacobian L{fj} logit L{fl} "
          f"{fl - fj} layers earlier")

```



English 'The capital of France is' jacobian L54 logit L55 1 layers earlier

Chinese '法国的首都是' jacobian L23 logit L58 35 layers earlier

French 'La capitale de la France est' jacobian L21 logit L57 36 layers earlier

Japanese 'フランスの首都は' jacobian L39 logit L56 17 layers earlier

三條非英文 prompt 都是同一個樣子,中文和法文的差距最大,Jacobian lens 在 L21 到 L23 就把 Paris 排進前五,logit lens 要等到 L57 以後。

英文問句用這個指標看不出差距(L54 對 L55),因為英文的 Paris 本來就不會晚進 top-5,要看機率才清楚,前面那張表的 L55 是 0.09 對 0.90。

三個語言各一條 prompt,推不到「所有非英文問句」。

5. J-space 到底是什麼

到這裡讀出的都是「哪個 token 分數最高」, J-space 是比這更強的宣稱。

詞表裡每個 token t 在第 ℓ 層的 residual 裡對應一個方向

$$v_t = (W_U J_\ell)_t$$

也就是 $W_U J_\ell$ 的第 t 列。讀出時那個 token 的分數就是內積 $\langle v_t, h \rangle$, 前面每一次 readout 都是在算所有 token 的這個內積。

J-space 說的是 residual 的一小塊可以寫成少數幾個這種方向的非負組合:

$$h_\ell \approx \sum_{t \in S} a_t v_t, \quad a_t \geq 0, \quad |S| \leq 25$$

S 是「此刻在 workspace 裡的東西」。非負的理由是這些方向代表「有這個概念」, 負係數沒有對應的意思。

詞表有 248320 個方向, 不可能全試, 所以先用 readout 分數挑前 k 個候選, 再解一個非負最小平方(投影梯度, 步長取 $1/\|G\|_F$ 保證不發散)。 k 掃 5、25、100, 另外拿隨機挑的 25 個 token 方向當底線。

```
In [16]: lh = llm.lm_head
```

```
def unembed_row(token):
    """One row of W_U, dequantized out of the 4-bit lm_head."""
    i = mx.array([token])
    return mx.dequantize(lh.weight[i], lh.scales[i], lh.biases[i],
                        group_size=lh.group_size,
                        bits=lh.bits)[0].astype(mx.float32)

def jlens_vec(token, l):
    """The J-lens direction of one token: row t of W_U @ J_l."""
    return unembed_row(token) @ Jlens[l].astype(mx.float32)

def nnls(V, h, iters=4000):
    """min ||h - V a|| with a >= 0, by projected gradient.

    The step is 1 / ||G||_F, and the Frobenius norm bounds the spectral
    norm from
    above, so the iteration cannot diverge (a bad step size showed up as
    negative
    explained variance the first time round).
    """
    G = V.T @ V
    b = V.T @ h
    L = float(mx.sqrt(mx.sum(G * G)))
    a = mx.zeros((V.shape[1],))
    for _ in range(iters):
        a = mx.maximum(a - (G @ a - b) / L, 0.0)
    return a

def jspace_fit(h, l, k=25, tokens=None):
    """Fit h with k J-lens directions: the k best readout scores unless
    given."""
    if tokens is None:
        sc = readout(transport(h[None, None], l))[0, 0].astype(mx.float32)
        tokens = [int(t) for t in mx.argsort(-sc)[:k].tolist()]
    V = mx.stack([jlens_vec(t, l) for t in tokens], axis=1)
    a = nnls(V, h)
    r = h - V @ a
    return 1.0 - float(mx.sum(r * r) / mx.sum(h * h)), tokens, a

ids = encode("The capital of France is")
res = residuals(ids)
last = ids.shape[1] - 1
KS = (5, 25, 100)
LAYERS = (48, 55, 60)
RANDOM_25 = [int(x) for x in mx.random.randint(0, 200000, (25,)),

key=mx.random.key(0)).tolist()]

ev, rnd, coef25, tok25 = {}, {}, None, None
for l in LAYERS:
    h = res[l][0, last].astype(mx.float32)
    for k in KS:
```



```

e, toks, a = jspace_fit(h, l, k)
ev[(l, k)] = e
if k == 25 and l == 60:
    coef25, tok25 = a, toks
if k == 25:
    ev[(l, "nz")] = int(mx.sum(a > 1e-6))
rnd[l] = jspace_fit(h, l, tokens=RANDOM_25)[0]


fig, axes = plt.subplots(1, 2, figsize=(7.4, 2.9),
                        gridspec_kw={"width_ratios": [1, 1.25]})

ax = axes[0]
for l, c in zip(LAYERS, (MUTED, JLENS, INK)):
    ax.plot(KS, [ev[(l, k)] * 100 for k in KS], marker="o", color=c)
    ax.text(KS[-1] * 1.15, ev[(l, KS[-1])] * 100, f"L{l}", color=c,
            fontsize=8,
            va="center")
ax.plot(KS, [rnd[LAYERS[-1]] * 100] * len(KS), color=LOGIT, ls="--",
        lw=1.0)
ax.text(KS[0], rnd[LAYERS[-1]] * 100 + 0.35, "random 25 directions",
        color=LOGIT,
        fontsize=7.5, va="bottom")
ax.set_xscale("log")
ax.set_xticks(KS)
ax.set_xticklabels([str(k) for k in KS])
ax.set_xlim(4.2, 260)
panel(ax, "how much of the residual it accounts for", "k (directions
used)",
      "explained %", 12)

ax = axes[1]
order = mx.argsort(-coef25).tolist()
vals = [float(coef25[i]) for i in order]
ax.bar(range(25), vals, color=[JLENS if v > 1e-6 else "#dfe3e4" for v in
vals],
      width=0.72)
shown = 0
for i in range(25):
    tk = tokenizer.decode([tok25[order[i]]]).strip()
    if shown >= 3 or not tk.isalpha():
        continue
    ax.text(i, vals[i] + max(vals) * 0.04, tk, rotation=90, fontsize=7,
           color=INK, ha="center", va="bottom")
    shown += 1
ax.set_ylim(0, max(vals) * 1.45)
panel(ax, "L60: the 25 coefficients, sorted", "candidate (sorted)",
      "coefficient", None)
plt.show()

for l in LAYERS:
    print(f"L{l} " + " ".join(f"k={k}: {ev[(l, k)] * 100:5.1f}%" for k
in KS)
          + f" random-25: {rnd[l] * 100:4.1f}% nonzero {ev[(l,
'nz')]} / 25")
print("L60 candidates, by readout score: "
      + " ".join(repr(tokenizer.decode([t])) for t in tok25[:6]))
print("L60 largest coefficients: "
      + " ".join(repr(tokenizer.decode([tok25[i]])) for i in order[:6]))

```



L48 $k=5$: 1.2% $k=25$: 1.5% $k=100$: 3.6% random-25: 0.2% nonzero 7/25

L55 $k=5$: 5.3% $k=25$: 6.7% $k=100$: 9.6% random-25: 0.1% nonzero 12/25

L60 $k=5$: 5.8% $k=25$: 8.1% $k=100$: 10.7% random-25: 0.0% nonzero 12/25

L60 candidates, by readout score: 'Paris' 'Paris' 'London' 'Berlin' 'paris' '巴黎'

L60 largest coefficients: 'Paris' '\n\n' '____' 'Berlin' 'Rome' 'Madrid'

25 個方向在 L48 解釋 1.5%, L55 是 6.7%, L60 是 8.1%, 放到 100 個方向最多也才 10.7%。residual 的絕大部分不在這些方向撈出來的東西裡面, 論文給的量級也是這樣, 所以不是實作沒做好。

隨機挑的 25 個方向只解釋 0.0% 到 0.2%, 差一到兩個數量級, 挑出來的那 25 個是有結構的。

非負最小平方還自己把一半的候選壓到 0, L55 和 L60 各只有 12 個係數是正的, L48 只有 7 個, 所以 $k = 25$ 是上限而不是實際用到的數目。

被挑進候選的是同一個概念的不同拼法加上同類的其他首都: L60 的前六名是 'Paris'、'Paris'、'London'、'Berlin'、'paris'、'巴黎'。但最後拿到正係數的那幾個不完全是這些, 係數最大的是 'Paris'、'Berlin'、'Rome'、'Madrid', 加上兩個排版碎片。也就是它只留一個 Paris 的拼法, 其餘的重量分給了同類的首都。

這裡也可以說明「J-space」這個名字的問題。非負加上稀疏上限, 這不是線性子空間, 不能講「投影到 J-space」; 照論文自己的 methods, 它是一堆多面錐的聯集。

6. 換掉 J-space 裡的概念

對齊原文的實驗:四個 prompt 分別問 France 的首都、語言、所屬洲、貨幣,在 J-space 裡把 France 換成 China, 四個情境用完全相同的介入。四個答案一起改的話,最簡單的解釋是它們讀的是同一份表徵。

做法是論文的 lens coordinate patching,用的就是 §5 那些 v_t :

$$V = [v_s \ v_t], \quad c = V^\dagger h, \quad h \leftarrow h + \alpha V (\sigma(c) - c)$$

c 是 h 在這兩個方向上的座標, σ 把兩個座標交換,跟 $\text{span}\{v_s, v_t\}$ 正交的成分不動。幾何上只動了這個平面;語義上是不是只換掉 France 這一個概念,公式不保證,要看輸出。

$V^\dagger$ 是 pseudo-inverse。不能直接拿內積當座標,因為 J-lens 向量 overcomplete 而且彼此不正交,下面會看到 v_{France} 跟 v_{China} 的 cosine 有 0.44, 直接取內積會把重疊的部分算兩次。

```
In [17]: def coords(x, vs, vt):
          V = mx.stack([vs, vt], axis=1)
          return V, mx.linalg.pinv(V, stream=mx.cpu) @ x

def swap(x, vs, vt, alpha):
    V, c = coords(x, vs, vt)
    return x + alpha * (V @ (c[::-1] - c))

def add_only(x, vs, vt, alpha):
    V, c = coords(x, vs, vt)
    return x + alpha * (c[0] - c[1]) * vt

FRANCE, CHINA = tokenizer.encode(" France")[0], tokenizer.encode(" China")
[0]
vf, vc = jlens_vec(FRANCE, 48), jlens_vec(CHINA, 48)
cos = float(vf @ vc / (mx.linalg.norm(vf) * mx.linalg.norm(vc)))
print(f"v_France at L48: {tuple(vf.shape)}, norm
{float(mx.linalg.norm(vf)):.1f}")
print(f"cos(v_France, v_China) = {cos:+.3f}")
```

v_France at L48: (5120,), norm 1.4

cos(v_France, v_China) = +0.441

三個選擇決定成不成:

- 層:L40 到 L55。層數是掃出來的,更早的層改了幾乎沒反應,更晚的層改了會把概念直接說出來(§7.5 有數字)。
- 位置:只改 France 那個 token,不碰最後一個位置。最後幾層的 J-lens 方向是「準備要說出口的字」,在那裡動手 model 會輸出 China 而不是拿它去推理。
- 強度: $\alpha = 1$ 是嚴格的交換, $\alpha = 2$ 是推過頭的版本,兩個都跑。

對照組 add_only 只加上 China 的成分,不把 France 拿掉。

```

In [18]: BAND = [l for l in Jlens if 40 <= l < 56]
ALPHA = 2.0

def patched(ids, pos, alpha, edit=swap):
    h = inner.embed_tokens(ids)
    fa, ssm, p = prep(h)
    for l, layer in enumerate(inner.layers):
        h = layer(h, mask=(ssm if layer.is_linear else fa),
                    cache=None, position_ids=p, position_embeddings=None)
    if alpha and l in BAND:
        h[0, pos] = edit(h[0, pos].astype(mx.float32),
                          jlens_vec(FRANCE, l), jlens_vec(CHINA, l),
                          alpha).astype(dtype)
    return readout(h)[0, ids.shape[1] - 1].astype(mx.float32)

def p_of(sc, s):
    return float(mx.softmax(sc)[tokenizer.encode(s)[0]])

def top1(sc):
    return tokenizer.decode([int(mx.argmax(sc))])

TASKS = [("capital", "The capital of France is", " Paris", " Beijing"),
          ("language", "The language spoken in France is", " French", " Chinese"),
          ("continent", "The continent that contains France is", " Europe", " Asia"),
          ("currency", "The currency of France is called the", " Euro", " Yuan")]
CONDS = ["before", "swap a=1", "swap a=2", "add only a=2"]

table = {}
for name, prompt, orig, new in TASKS:
    ids = encode(prompt)
    toks = [tokenizer.decode([int(t)]) for t in ids[0].tolist()]
    pos = next(i for i, t in enumerate(toks) if "France" in t)
    for cond, sc in zip(CONDS, [patched(ids, pos, 0),
                                patched(ids, pos, 1.0),
                                patched(ids, pos, ALPHA),
                                patched(ids, pos, ALPHA, add_only)]):
        table[(name, cond)] = (p_of(sc, orig), p_of(sc, new), top1(sc))

fig, axes = plt.subplots(1, 4, figsize=(7.6, 2.6), sharey=True)
x = mx.arange(len(CONDS)).tolist()
for ax, (name, _, orig, new) in zip(axes, TASKS):
    o = [table[(name, c)][0] for c in CONDS]
    n = [table[(name, c)][1] for c in CONDS]
    ax.bar([i - 0.2 for i in x], o, width=0.38, color=LOGIT)
    ax.bar([i + 0.2 for i in x], n, width=0.38, color=SWAPPED)
    ax.set_xticks(x)
    ax.set_xticklabels(["before", "a=1", "a=2", "add\nonly"],
                        fontsize=7.5)
    panel(ax, name, f"{orig.strip()} -> {new.strip()}", "", 0.72)
axes[0].set_ylabel("probability")
axes[3].text(0.97, 0.95, "original answer", color=LOGIT, fontsize=8,
             transform=axes[3].transAxes, va="top", ha="right")


```

```

axes[3].text(0.97, 0.82, "swapped-in", color=SWAPPED, fontsize=8,
             transform=axes[3].transAxes, va="top", ha="right")
plt.show()

for name, prompt, orig, new in TASKS:
    print(f"{prompt!r}")
    for c in CONDS:
        po, pn, t = table[(name, c)]
        print(f"    {c:13s} {t!r:11s} p({orig})={po:.3f} p({new})=
{pn:.3f}")

```



'The capital of France is'

before	' Paris'	p(Paris)=0.617	p(Beijing)=0.000
swap a=1	' Paris'	p(Paris)=0.497	p(Beijing)=0.034
swap a=2	' Beijing'	p(Paris)=0.006	p(Beijing)=0.383
add only a=2	' Paris'	p(Paris)=0.427	p(Beijing)=0.066

'The language spoken in France is'

before	' French'	p(French)=0.539	p(Chinese)=0.000
swap a=1	' French'	p(French)=0.528	p(Chinese)=0.001
swap a=2	' Chinese'	p(French)=0.004	p(Chinese)=0.222
add only a=2	' French'	p(French)=0.461	p(Chinese)=0.002

'The continent that contains France is'

before	' Europe'	p(Europe)=0.153	p(Asia)=0.001
swap a=1	' Europe'	p(Europe)=0.135	p(Asia)=0.007
swap a=2	' Asia'	p(Europe)=0.090	p(Asia)=0.096
add only a=2	' Europe'	p(Europe)=0.138	p(Asia)=0.005

'The currency of France is called the'

before	' euro'	p(Euro)=0.111	p(Yuan)=0.000
swap a=1	' euro'	p(Euro)=0.101	p(Yuan)=0.017
swap a=2	' yuan'	p(Euro)=0.001	p(Yuan)=0.060
add only a=2	' euro'	p(Euro)=0.097	p(Yuan)=0.036

四個問題用同一個介入,答案全部跟著改:Paris 變 Beijing、French 變 Chinese、Europe 變 Asia、euro 變 yuan,四個都需要 $\alpha = 2$ 。

$\alpha = 1$ 的嚴格交換方向對但推不過去, $p(\text{Beijing})$ 從 0.000 到 0.034, top-1 還是 'Paris', 四個任務一個都沒翻。對照組 add only $\alpha=2$ 四個也都沒翻, $p(\text{Beijing})$ 只有 0.066, 交換版是 0.383, 所以「把 France 拿掉」這一半是必要的。

幅度要留意。Neel Nanda 團隊在同一顆 model 上複現時也只拿到微弱但為正的因果效果,這是一顆 27B 的 4-bit 量化 model,不是 Sonnet 4.5。Asia 那一列最勉強, 0.096 對 Europe 的 0.090, 剛好過去而已。

7. 論文的五個 workspace 面向

論文給了五個功能面向,要一項一項成立才算 workspace (原文圖 2):

面向	問題
verbal report	問它在想什麼,它說得出來嗎
directed modulation	叫它想 X、或叫它不要想 X, J-space 跟著動嗎
internal reasoning	它會拿 J-space 裡的東西去推理,還是只是存著
flexible generalization	同一個概念,換到不同任務還通嗎
selectivity	只有進 workspace 的內容才影響行為嗎

flexible generalization 就是 §6 做的事,一個介入、四個任務一起改。這一節跑剩下四個, 另外加一個論文用來驗因果的手法:把概念消除。

介入先寫成一個通用的 forward, edits 是一串 (position, 函式), 每一層都套上去。

```

In [19]: SCAN = [l for l in JlenS if 40 <= l < 62 and l % 2 == 0]

def tid(word):
    """First token id of a spelling."""
    return tokenizer.encode(word)[0]

def edited(ids, edits=(), layers=BAND):
    """Forward pass with edits = [(position, fn)] applied at every layer in
    `layers`."""
    h = inner.embed_tokens(ids)
    fa, ssm, p = prep(h)
    for l, layer in enumerate(inner.layers):
        h = layer(h, mask=(ssm if layer.is_linear else fa),
                  cache=None, position_ids=p, position_embeddings=None)
        if l in layers:
            for pos, fn in edits:
                h[0, pos] = fn(h[0, pos].astype(mx.float32),
                               l).astype(dtype)
    return readout(h)[0, ids.shape[1] - 1].astype(mx.float32)

def best_readout(ids, words, layers=SCAN):
    """Largest J-lens probability of any spelling, over every layer and
    position."""
    res = residuals(ids)
    best = (0.0, None, None)
    for l in layers:
        sm = mx.softmax(readout(transport(res[l], l)).astype(mx.float32),
                        axis=-1)[0]
        for w in words:
            col = sm[:, tid(w)]
            p = float(mx.max(col))
            if p > best[0]:
                best = (p, l, int(mx.argmax(col)))
    return best

def probs(sc, *pairs):
    sm = mx.softmax(sc)
    return " ".join(f"p({n})={max(float(sm[tid(w)]) for w in ws):.3f}"
                    for n, ws in pairs)

def where(ids, word):
    toks = [tokenizer.decode([int(t)]) for t in ids[0].tolist()]
    return next(i for i, t in enumerate(toks) if word in t)

print(f"scan layers: {SCAN[0]}..{SCAN[-1]} ({len(SCAN)})    patch band: "
      f"L{BAND[0]}-L{BAND[-1]}")

```

scan layers: 40..60 (11) patch band: L40-55

7.1 定向調節:叫它想,叫它不要想

同一句抄寫任務,後面接不同的指令,量 J-space 裡有沒有那個概念。兩組設計:

- 白熊。 `bear` 在兩個 prompt 裡都出現,「因為輸入有這個字才讀到」的效應互相抵銷,差的只有指令。
- lemon。指令說的是 citrus fruit, `lemon` 從頭到尾沒出現在 prompt 裡,讀到它就不能用 input-copying 解釋。對照組把指令換成想呼吸,測「任何指令都會抬高一切」。

讀法是掃 L40 到 L60 每一層、每一個位置,取最大的機率。


```
In [20]: S = 'Write the sentence "The old painting hung crookedly on the wall."'

CASES = [
    ("bear (in the prompt)", ("bear", "bears"), [
        ("not mentioned", S),
        ("think about it", S + " While you write, think about a white bear."),
        ("do NOT think of it", S + " While you write, do not think about a white bear.")] ),
    ("lemon (never in the prompt)", ("lemon", "lemons"), [
        ("no instruction", S),
        ("think about citrus", S + " While you write, think about citrus fruit."),
        ("think of breathing", S + " While you write, think about your breathing.")] ),
]

got = []
for target, words, cases in CASES:
    for tag, prompt in cases:
        p, l, pos = best_readout(encode(prompt), words)
        got.append((target, tag, p, l, pos))

fig, axes = plt.subplots(1, 2, figsize=(7.2, 2.9), sharey=True)
for ax, (target, words, cases) in zip(axes, CASES):
    rows = [g for g in got if g[0] == target]
    vals = [g[2] for g in rows]
    ax.bar(range(len(rows)), vals,
           color=[LOGIT, JLENS, SWAPPED][:len(rows)], width=0.62)
    for i, v in enumerate(vals):
        ax.text(i, v * 1.25, f"{v:.4f}", ha="center", fontsize=7.5, color=INK)
    ax.set_xticks(range(len(rows)))
    ax.set_xticklabels([g[1].replace(" ", "\n") for g in rows],
                       fontsize=7.5)
    ax.set_yscale("log")
    ax.set_ylim(5e-5, 1.0)
    panel(ax, target, "", "", None)
axes[0].set_ylabel("max J-lens p over layers x positions")
plt.show()

for target, tag, p, l, pos in got:
    print(f"{target:28s} {tag:20s} max p = {p:.4f} at L{l} pos{pos}")
```

bear (in the prompt)	not mentioned	max p = 0.0001	at
L60 pos5			
bear (in the prompt)	think about it	max p = 0.0075	at
L54 pos23			
bear (in the prompt)	do NOT think of it	max p = 0.1079	at
L60 pos24			
lemon (never in the prompt)	no instruction	max p = 0.0001	at
L48 pos3			
lemon (never in the prompt)	think about citrus	max p = 0.0114	at
L58 pos22			
lemon (never in the prompt)	think of breathing	max p = 0.0001	at
L48 pos3			

白熊那組成立,而且方向是反的。沒提到白熊時 J-space 裡幾乎沒有它, max p = 0.0001;叫它想白熊升到 0.0075;叫它不要想白熊升到 0.1079,比叫它想的時候 高一個數量級以上。抑制的指令沒有把概念壓下去,反而把它推進 workspace, 論文圖 10 有跨 model 的版本。

lemon 那組是乾淨的正向。指令說的是 citrus fruit, **lemon** 從沒出現在 prompt 裡, J-space 讀到 0.0114;把指令換成想呼吸掉回 0.0001,跟完全沒有指令一樣。動的是被指定的那個概念,不是有指令就整體抬高。

絕對值都很小,0.0001 到 0.1,要看的是同一個 target 在不同指令之間的比值。

7.2 把概念消除

交換是把座標換過去,消除是把座標歸零,同一套機制取 $\sigma(c) = 0$:

$$V = [v_1 \cdots v_k], \quad h \leftarrow h - \alpha V (V^\dagger h)$$

一次投掉整組方向。J-lens 向量彼此不正交,一個一個減會重複扣掉共用的成分。

介入從最小的開始加碼:一個方向、一個位置,到一個方向、所有位置,再到整組 {France, French, Paris}、所有位置。另外也試直接把答案方向 **Paris** 拿掉。

```

In [21]: def ablate_span(words, alpha=1.0):
          """h <- h - alpha * V (V+ h): project out the span of several J-lens
          directions."""
          def f(x, l):
              V = mx.stack([jlens_vec(tid(w), l) for w in words], axis=1)
              return x - alpha * (V @ (mx.linalg.pinv(V, stream=mx.cpu) @ x))
          return f

          PROMPT = "The capital of France is"
          ids = encode(PROMPT)
          fpos, last = where(ids, "France"), ids.shape[1] - 1
          WIDE = [l for l in Jlen if 20 <= l < 62]
          ALL = list(range(ids.shape[1]))


          RUNS = [
              ("baseline", (), BAND),
              ("France @France, L40-55", ((fpos, ablate_span([" France"])),),
              BAND),
              ("France @France, L20-61", ((fpos, ablate_span([" France"])),),
              WIDE),
              ("France @every pos, L40-55", tuple((p, ablate_span([" France"])) for
              p in ALL), BAND),
              ("{France,French,Paris} @every pos", tuple((p, ablate_span(["
              France", " French", " Paris"]))) for p in ALL), BAND),
              ("Paris @last pos, L40-55", ((last, ablate_span([" Paris"])),),
              BAND),
              ("Paris @last pos, L20-61", ((last, ablate_span([" Paris"])),),
              WIDE),
              ]

          out = []
          for tag, edits, layers in RUNS:
              sc = edited(ids, edits, layers)
              out.append((tag, float(mx.softmax(sc)[tid(" Paris")]), top1(sc)))

          fig, ax = plt.subplots(figsize=(6.2, 2.9))
          y = list(range(len(out))[:-1])
          vals = [o[1] for o in out]
          ax.barh(y, vals, color=[LOGIT if o[2] == " Paris" else SWAPPED for o in
          out],
                  height=0.66)
          for yi, (tag, v, t) in zip(y, out):
              ax.text(v + 0.012, yi, f"{v:.3f}" + (" if t == " Paris" else f" top-
              1 {t!r}"),
                      va="center", fontsize=7.5, color=INK)
          ax.set_yticks(y)
          ax.set_yticklabels([o[0] for o in out], fontsize=7.5)
          ax.set_xlim(0, 0.82)
          ax.grid(axis="x")
          ax.set_axisbelow(True)
          ax.set_title(f"p(Paris) for {PROMPT!r}", loc="left", color=INK)
          ax.set_xlabel("p(Paris)")
          plt.show()

          for tag, v, t in out:
              print(f"{tag:34s} p(Paris)={v:.3f}    top1={t!r}")

```



baseline	$p(\text{Paris})=0.617$	top1=' Paris '
-France @France, L40-55	$p(\text{Paris})=0.524$	top1=' Paris '
-France @France, L20-61	$p(\text{Paris})=0.559$	top1=' Paris '
-France @every pos, L40-55	$p(\text{Paris})=0.197$	top1=' Paris '
-{France, French, Paris} @every pos	$p(\text{Paris})=0.001$	top1=' not '
-Paris @last pos, L40-55	$p(\text{Paris})=0.040$	top1=' not '
-Paris @last pos, L20-61	$p(\text{Paris})=0.000$	top1=' not '

消除成不成看你動多少。

- 投掉 France 一個方向,只在 France 那個位置: $p(\text{Paris})$ 從 0.617 到 0.524。
- 同一個方向,層數放寬到 L20-61:0.559。
- 一個方向,每個位置都投:0.197,掉了三分之二。
- 整組 {France, French, Paris}、每個位置:0.001,top-1 變成 ' not '。
- 直接投掉答案方向 Paris,只在最後一個位置:0.040;層數放寬到 L20-61 是 0.000。

概念不是只存在提到它的那個位置上, France 的資訊在 forward 的過程中已經散到別的位置,只掐一個點沒用。要投也得投整組,一個一個減會漏掉共用的成分。

7.3 選擇性:明確報告 vs 自動處理

論文說進了 workspace 的內容只在需要報告或彈性使用時起作用,自動處理那條路走不到它,原文圖 8 的例子是一段西班牙文。

兩個 prompt 共用同一句西班牙文,介入完全相同,把西班牙文那幾個 token 的 Spanish 座標換成 French:

- 明確報告:問這句是什麼語言。翻得動代表這個概念被拿去回答問題。
- 自動處理:叫它照原樣接下去。論文的說法成立的話,這條不該受影響,照樣接西班牙文。

強度要掃過去,不能挑一個 α 交差。

```

In [22]: def swap_lang(alpha, src=" Spanish", tgt=" French"):
          return lambda x, l: swap(x, jlens_vec(tid(src), l),
          jlens_vec(tid(tgt), l), alpha)


EXPLICIT = 'Sentence: "El sol se pone en el oeste." That sentence is
written in'
AUTO = 'Continue in the same language: "El sol se pone en el'
ALPHAS = (0, 0.5, 1.0, 1.5, 2.0, 3.0)

runs = []
for prompt, keys, lastw in [
    (EXPLICIT, (("Spanish", " Spanish"), ("French", " French")), "
oeste"),
    (AUTO, (("oeste", " oeste"), ("ouest", " ouest")), " el")]:
    ids = encode(prompt)
    toks = [tokenizer.decode([int(t)]) for t in ids[0].tolist()]
    a = where(ids, "El")
    b = max(i for i, t in enumerate(toks) if t == lastw)
    ps = list(range(a, b + 1))
    series = {name: [] for name, _ in keys}
    tops = []
    for alpha in ALPHAS:
        edits = () if alpha == 0 else tuple((p, swap_lang(alpha)) for p in
ps)
        sc = edited(ids, edits)
        for name, w in keys:
            series[name].append(float(mx.softmax(sc)[tid(w)]))
        tops.append(top1(sc))
    runs.append((prompt, ps, series, tops))

fig, axes = plt.subplots(1, 2, figsize=(7.4, 2.9))
TITLES = ["explicit report: what language is this?",
          "automatic processing: continue the sentence"]
for ax, (prompt, ps, series, tops), title in zip(axes, runs, TITLES):
    for (name, vals), c in zip(series.items(), (LOGIT, SWAPPED)):
        ax.plot(ALPHAS, vals, marker="o", ms=3.5, color=c)
        if max(vals) > 0.002:
            j = vals.index(max(vals))
            ax.annotate(name, (ALPHAS[j], vals[j]), textcoords="offset
points",
                        xytext=(5, 4), color=c, fontsize=8)
        ax.axvline(1.5, color=MUTED, lw=0.9, ls=":")
        ax.set_xlim(-0.15, 3.9)
        top = max(max(v) for v in series.values())
        panel(ax, title, "alpha", "", top * 1.35 + 0.005)
axes[0].set_ylabel("probability")
axes[1].set_ylabel("probability")
axes[1].text(1.6, 0.02, "continuation\nbreaks here", color=MUTED,
fontsize=7.5,
            va="top")
plt.show()

for prompt, ps, series, tops in runs:
    print(f"{prompt!r}    patched positions {ps[0]}..{ps[-1]}")
    for i, alpha in enumerate(ALPHAS):
        cols = " ".join(f"p({n})={v[i]:.3f}" for n, v in series.items())
        print(f"    alpha={alpha:<4} {cols}    top1={tops[i]!r}")

```



'Sentence: "El sol se pone en el oeste." That sentence is written in' patched positions 3..9

alpha=0	p(Spanish)=0.108	p(French)=0.006	top1=' which'
alpha=0.5	p(Spanish)=0.103	p(French)=0.014	top1=' which'
alpha=1.0	p(Spanish)=0.099	p(French)=0.015	top1=' which'
alpha=1.5	p(Spanish)=0.003	p(French)=0.131	top1=' which'
alpha=2.0	p(Spanish)=0.002	p(French)=0.123	top1=' which'
alpha=3.0	p(Spanish)=0.002	p(French)=0.119	top1=' which'

'Continue in the same language: "El sol se pone en el' patched positions 7..12

alpha=0	p(oeste)=0.014	p(ouest)=0.000	top1=' horizonte'
alpha=0.5	p(oeste)=0.015	p(ouest)=0.000	top1=' horizonte'
alpha=1.0	p(oeste)=0.017	p(ouest)=0.000	top1=' horizonte'
alpha=1.5	p(oeste)=0.000	p(ouest)=0.000	top1=' Spanish'
alpha=2.0	p(oeste)=0.000	p(ouest)=0.000	top1=' Spanish'
alpha=3.0	p(oeste)=0.000	p(ouest)=0.000	top1=' Spanish'

只成立一半。

明確報告那條翻得動,但要 $\alpha = 1.5$ 。 $\alpha \leq 1.0$ 時 p(Spanish) 還是 0.099、
p(French) 0.015,到 $\alpha = 1.5$ 時 p(Spanish) 掉到 0.003, p(French) 升到 0.131。

自動接續那條在 $\alpha \leq 1.0$ 完全不受影響,照樣接 ' horizonte', p(oeste) 甚至 微升。
到 $\alpha = 1.5$ 就壞了,而且不是變成法文,top-1 變成 ' Spanish', 它開始講語言的名字而不是接句子。

兩條路的門檻是同一個,能讓明確報告翻過去的強度剛好也是讓自動接續脫軌的強度, 所以拿不到論文那個乾淨的分離。能說的只有:在自動接續還正常的強度下,概念的改動 沒有影響它。

7.4 口頭報告:注入一個概念,問它在想什麼

把一個概念寫進 residual stream,再問它腦裡有什麼,論文的結果是它會說出那個字。

「它說出那個字」有兩種解釋:它讀到了自己 workspace 裡的內容再報告出來,或者注入直接偏移了輸出分布,不管問什麼都會說。所以要有對照組,同一個注入換一個完全無關的問題,如果那邊也照樣說出來就不能算報告。

```
In [23]: ASK = "What single word is on your mind right now? Answer with one word:"
CONTROL = "The capital of France is"

def inject(word, alpha):
    def f(x, l):
        v = jlens_vec(tid(word), l)
        return x + alpha * v / mx.linalg.norm(v) * mx.linalg.norm(x)
    return f

def generate(ids, n=5, edits=(), layers=BAND):
    """Greedy continuation, re-running the whole forward for each token."""
    cur, out = ids, []
    for _ in range(n):
        t = int(mx.argmax(edited(cur, edits, layers)))
        if t in tokenizer.all_special_ids:
            break
        out.append(t)
        cur = mx.concatenate([cur, mx.array([[t]])], axis=1)
    return tokenizer.decode(out)

for word, alpha in (("volcano", 0.0), ("volcano", 0.5), ("volcano", 1.0),
                    ("guitar", 0.5)):
    line = []
    for prompt, n in ((ASK, 5), (CONTROL, 4)):
        ids = encode(prompt)
        edits = () if alpha == 0 else tuple((p, inject(word, alpha))
                                             for p in range(ids.shape[1]))
        line.append(generate(ids, n, edits))
    print(f"inject{word:9s} alpha={alpha:<4} ask -> {line[0]!r:26s}"
          f" control -> {line[1]!r}")
```

```
inject volcano alpha=0.0 ask -> "\n\n<think>\nHere's"
control -> ' Paris.\n\n<think>'
inject volcano alpha=0.5 ask -> ' volcano.\n\n<think>\n'
control -> ' volcano.\n\n<think>'
inject volcano alpha=1.0 ask -> ' volcano.\n\n<think>\n'
control -> ' volcano, and the'
inject guitar alpha=0.5 ask -> ' guitar.\n\n<think>\n'
control -> ' guitar.\n\n<think>'
```

不成立,對照組把它打掉了。

注入之後它確實會說。問「你現在想到哪一個字」, $\alpha = 0.5$ 答 ' volcano.', 換成 guitar 答 ' guitar.', 沒注入的時候它只會進 <think>。

但對照組也照樣說。同一個注入問 The capital of France is ,答的是 volcano 或 guitar 而不是 Paris, α 從 0.5 到 1.0 都一樣。

在這顆 model 上,這個注入只是把輸出分布推向那個字,不能算報告 workspace 的內容。要分開這兩件事,得找到一個強度或位置讓「問它想什麼」會說出來、「問別的問題」不會,我掃過的範圍裡沒有。

7.5 內部推理:交換之後,它會拿新概念去算嗎

這一項要求交換之後 model 用新概念做推論,而不是把它輸出。兩種題型。

數腳,論文圖 6 的例子:把 spider 換成 ant,問腳的數量,答案該從 8 變 6。8 和 6 都不在 prompt 裡,也不是被交換的那兩個字。這顆 model 是 thinking model,問句形式會讓它先吐 \n\n,所以用完成句的寫法。

兩跳事實: The capital of the country where the Eiffel Tower stands is 答 Paris, 中間那一跳 France 完全不在 prompt 裡,是 model 自己算出來的。把 France 換成 China 之後, 答 Beijing 代表它拿新概念又跑了一次第二跳,答 China 代表它只是把注入的東西講出來。這個題型能把推理和洩漏分開。

第二種還要挑位置。France 不在輸入裡,沒有「France 那個 token 的位置」可以改,只能改別的位置,三個選擇都試:最後一個位置、Tower 那個位置、除了最後一個以外全改。

```

In [24]: EIGHT, SIX = ("8", " eight"), ("6", " six")
ALPHAS = (0, 1.0, 2.0, 3.0)

legs = {}
for prompt, src, tgt, keep, want in [
    ("A spider has", " spider", " ant", EIGHT, SIX),
    ("The number of legs on a spider is", " spider", " ant", EIGHT,
SIX)]:
    ids = encode(prompt)
    pos = where(ids, src.strip())
    for alpha in ALPHAS:
        edits = () if alpha == 0 else (
            (pos, lambda x, l, a=alpha: swap(x, jlens_vec(tid(src), l),
                jlens_vec(tid(tgt), l), a)),)

        sc = edited(ids, edits)
        sm = mx.softmax(sc)
        legs[(prompt, alpha)] = (max(float(sm[tid(w)]) for w in keep),
                                max(float(sm[tid(w)]) for w in want),
                                top1(sc))

TWO_HOP = "The capital of the country where the Eiffel Tower stands is"
ids = encode(TWO_HOP)
last = ids.shape[1] - 1
TARGETS = [("last position", [last]),
            ("Tower position", [where(ids, "Tower")]),
            ("all but last", list(range(1, last)))]
WATCH = [("Paris", " Paris"), ("Beijing", " Beijing"), ("France", "
France")]

def swap_fc(alpha):
    return lambda x, l: swap(x, jlens_vec(tid(" France"), l),
                             jlens_vec(tid(" China"), l), alpha)

hop = {}
for tag, ps in TARGETS:
    for alpha in ALPHAS:
        edits = () if alpha == 0 else tuple((p, swap_fc(alpha)) for p in
ps)

        sc = edited(ids, edits)
        sm = mx.softmax(sc)
        hop[(tag, alpha)] = ({n: float(sm[tid(w)]) for n, w in WATCH},
top1(sc))

fig, axes = plt.subplots(1, 4, figsize=(8.4, 2.6))

ax = axes[0]
p = "A spider has"
ax.plot(ALPHAS, [legs[(p, a)][0] for a in ALPHAS], marker="o", ms=3.5,
color=LOGIT)
ax.plot(ALPHAS, [legs[(p, a)][1] for a in ALPHAS], marker="o", ms=3.5,
color=SWAPPED)
ax.text(0.97, 0.95, "p(8), spider's answer", color=LOGIT, fontsize=7.5,
        transform=ax.transAxes, va="top", ha="right")
ax.text(0.97, 0.82, "p(6), ant's answer", color=SWAPPED, fontsize=7.5,
        transform=ax.transAxes, va="top", ha="right")
panel(ax, "legs: 'A spider has'", "alpha", "probability", 0.18)
axes[1].set_ylabel("probability")

```




```

for ax, (tag, _) in zip(axes[1:], TARGETS):
    for (name, _), c in zip(WATCH, (LOGIT, SWAPPED, MUTED)):
        ys = [hop[(tag, a)][0][name] for a in ALPHAS]
        ax.plot(ALPHAS, ys, marker="o", ms=3.5, color=c)
        if max(ys) > 0.05:
            j = ys.index(max(ys))
            ax.annotate(name, (ALPHAS[j], ys[j]), textcoords="offset
points",
                        xytext=(5, 3), color=c, fontsize=7.5)
        ax.set_xlim(-0.15, 4.6)
        panel(ax, f"patch {tag}", "alpha", "", 1.05)
        ax.set_title(f"patch {tag}", loc="left", color=INK, fontsize=8.5)
plt.show()

for prompt in ("A spider has", "The number of legs on a spider is"):
    print(f"{prompt!r}    swap spider -> ant")
    for a in ALPHAS:
        k, w, t = legs[(prompt, a)]
        print(f"    alpha={a:<4} p(8)={k:.3f} p(6)={w:.3f} top1={t!r}")
print(f"\ntwo-hop: {TWO_HOP!r}")
print(f" 'France' appears in the prompt: {'France' in TWO_HOP}")
for tag, ps in TARGETS:
    print(f" patch {tag} {ps}")
    for a in ALPHAS:
        d, t = hop[(tag, a)]
        cols = " ".join(f"p({n})={d[n]:.3f}" for n, _ in WATCH)
        print(f"    alpha={a:<4} {cols} top1={t!r}")

```



'A spider has' swap spider -> ant

alpha	p(8)	p(6)	top1
0	0.124	0.004	' '
1.0	0.108	0.004	' '
2.0	0.092	0.003	' '
3.0	0.083	0.003	' '

'The number of legs on a spider is' swap spider -> ant

alpha	p(8)	p(6)	top1
0	0.050	0.009	' '
1.0	0.039	0.009	' '
2.0	0.008	0.007	' '
3.0	0.008	0.007	' '

two-hop: 'The capital of the country where the Eiffel Tower stands is'

'France' appears in the prompt: *False*

```

patch last position [11]
  alpha=0    p(Paris)=0.394  p(Beijing)=0.000  p(France)=0.001
top1=' Paris'
  alpha=1.0  p(Paris)=0.262  p(Beijing)=0.003  p(France)=0.000
top1=' Paris'
  alpha=2.0  p(Paris)=0.000  p(Beijing)=0.000  p(France)=0.975
top1=' France'
  alpha=3.0  p(Paris)=0.000  p(Beijing)=0.000  p(France)=0.980
top1=' France'
patch Tower position [9]
  alpha=0    p(Paris)=0.394  p(Beijing)=0.000  p(France)=0.001
top1=' Paris'
  alpha=1.0  p(Paris)=0.373  p(Beijing)=0.000  p(France)=0.001
top1=' Paris'
  alpha=2.0  p(Paris)=0.148  p(Beijing)=0.011  p(France)=0.000
top1=' Paris'
  alpha=3.0  p(Paris)=0.163  p(Beijing)=0.010  p(France)=0.000
top1=' Paris'
patch all but last [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
  alpha=0    p(Paris)=0.394  p(Beijing)=0.000  p(France)=0.001
top1=' Paris'
  alpha=1.0  p(Paris)=0.350  p(Beijing)=0.000  p(France)=0.001
top1=' Paris'
  alpha=2.0  p(Paris)=0.000  p(Beijing)=0.189  p(France)=0.000
top1=' Beijing'
  alpha=3.0  p(Paris)=0.000  p(Beijing)=0.218  p(France)=0.000
top1=' Beijing'

```

數腳那組不成立, $p(8)$ 從 0.124 掉到 0.083, $p(6)$ 從 0.004 完全沒動。交換把舊概念弄壞了,沒把新概念的後果裝上去。

兩跳事實那組成立,但位置決定一切:

- 只改最後一個位置: $\alpha = 2$ 時 $p(\text{France})$ 衝到 0.975, top-1 是 ' France '。它把被交換的概念直接說出來,這就是前面提的 late-layer 洩漏,現在有數字。
- 只改 Tower 那個位置:幾乎沒反應, $p(\text{Beijing})$ 最多 0.011,單一個位置扛不起來。
- 除了最後一個以外全改: $\alpha = 2$ 時 top-1 變成 ' Beijing ', $p(\text{Paris})$ 掉到 0.000, 而 $p(\text{France})$ 、 $p(\text{China})$ 都還在 0.000 到 0.001。

最後那個是要找的結果。它沒有把 China 講出來,講的是第二跳的結果,而被換掉的概念 從頭到尾不在輸入裡,所以也不可能是 input-copying。

兩個題型合起來,結論跟 Neel Nanda 團隊的複現一致:多跳事實的替換站得住,需要算的 那種不行。數腳要的是「知道 ant 有 6 隻腳」這種屬性查詢,兩跳事實要的是「拿 China 再跑一次首都查詢」,後者這顆 model 做得到。

六項跑完的樣子

面向	這顆 model(Qwen3.6-27B-4bit)
flexible generalization	成立。一個介入,四個任務一起改(\$6)
directed modulation	成立。叫它不要想白熊,白熊反而更清楚
internal reasoning	成立,但只有兩跳事實那種,數腳那種不行
概念消除(論文驗因果的手法)	成立,但要投整組方向、每個位置都投
selectivity	一半。明確報告翻得動,但拿不到跟自動處理的乾淨分離
verbal report	不成立。注入之後,問什麼都會說出那個字

四項成立、一項一半、一項不成立。這是一顆 27B 的 4-bit 量化 model,不是 Sonnet 4.5, 效應本來就該弱得多。

成立的那幾項有一個共同的條件:介入要落在內容位置,而且往往要落在一整片位置上。只改一個點不夠,改最後一個位置會退化成把概念直接說出來。

小結

- residual stream 是所有層共用的座標系,所以才有得讀。
- logit lens 把最後那顆頭借來讀中間層,便宜,但暗中假設 $J_\ell = I$ 。
- Jacobian lens 把 J_ℓ 算出來,讀同一條 residual 卻早好幾層看到答案,非英文問句的差距最大,中文和法文早了 30 幾層。
- 從零算 J_ℓ 只需要 `mx.vjp` 加兩個實務 workaround,一條 prompt 的結果就跟一千條 擬合的公開 lens 大致同向。
- J-space 是稀疏非負組合,25 個方向只解釋 residual 的 1.5% 到 8%,但比隨機挑的高 一到兩個數量級 (§5)。
- 改得動:同一個介入把 `France` 換成 `China`,四個不同的問題一起改口;把整組方向投掉,答案就消失。
- 五個面向裡四項成立、一項一半、一項不成立,而且成不成很吃介入落在哪些位置 (§7)。

該保留的懷疑

- 線性近似。 J_ℓ 是一階展開,後面幾十層並不是線性的。
- 全域先驗。公開 lens 是跨一千條 prompt 平均的,會帶進跟當前輸入無關的偏好。
- input-copying。如果要探測的字本來就出現在 prompt 裡,lens 會在那個位置直接把它讀成第一名,反映的是輸入而不是內部想法。上面的例子沒有踩到這點(`Paris` 不在 prompt 裡, §7.5 的 `France` 也不在),但要自己延伸實驗時得檢查 (issue #5)。
- 名字取得不好。J-space 不是 residual stream 的線性子空間,照論文自己的 methods,它是在稀疏限制下的一堆多面錐的聯集,§5 量過它只佔 1.5% 到 8%。
- 循環論證的疑慮。一個依「對最終輸出的影響」定義出來的空間,本來就會跟最終輸出高度相關。
- 不是每個實驗都能重現。獨立重現裡多步事實的替換站得住,詩韻規劃和心算沒有重現出來 (Nanda 團隊的 review、[tao-hpu/jspace-replication](#))。

參考資料

- [anthropics/jacobian-lens](#) 官方實作
- [neuronpedia/jacobian-lens](#) 擬合好的 lens
- [Interpreting GPT: the logit lens](#)
- [A Mathematical Framework for Transformer Circuits](#)
- [On the Biology of a Large Language Model](#)
- [J-Space: Yet Another LLM Mind Reader?](#)
- [A Review of Anthropic's Global Workspace Paper](#)
- [Gated Delta Networks](#)